

Linear Regression- Auto Dataset

Import Libraries

In [1]:

```
1
2 # common libraries
3 import numpy as np
4 import pandas as pd
5
6 # Visualization
7 import seaborn as sns
8 import matplotlib.pyplot as plt
9
10 # Feature selection
11 from statsmodels.stats.outliers_influence import variance_inflation_factor
12
13 # model building
14 from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV
15 from sklearn.linear_model import LinearRegression, Ridge, Lasso
16
17 # Evaluation Matrix
18 from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
19
20 # To avoid warning
21
22 import warnings
23 warnings.filterwarnings("ignore")
24
25 # To handle outliers
26 from scipy.stats import boxcox
27
28 # To check Normality of Residual
29 from scipy.stats import shapiro, kstest, normaltest
30 import statsmodels.api as sm
31
32 # To create pickle file
33 import pickle
34
35 # To create json file
36 import json
```

Step 1: Problem Statement

- | | |
|---|--|
| 1 | To predict price of the cars by using supervised machine learning considering as a regression problem. |
|---|--|


Step 2: Data Gathering

In [145]:

```
1 df = pd.read_csv("autos_dataset.csv")
2 df
```

Out[145]:

make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore
alfa-nero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47
alfa-nero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47
alfa-nero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68
audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19
audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19
...
volvo	gas	std	four	sedan	rwd	front	109.1	...	141	mpfi	3.78
volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78
volvo	gas	std	four	sedan	rwd	front	109.1	...	173	mpfi	3.58
volvo	diesel	turbo	four	sedan	rwd	front	109.1	...	145	idi	3.01
volvo	gas	turbo	four	sedan	rwd	front	109.1	...	141	mpfi	3.78



Step 3: Exploratory Data Analysis

In [3]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      205 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   aspiration              205 non-null    object
5   num-of-doors           205 non-null    object
6   body-style             205 non-null    object
7   drive-wheels           205 non-null    object
8   engine-location        205 non-null    object
9   wheel-base             205 non-null    float64
10  length                 205 non-null    float64
11  width                  205 non-null    float64
12  height                 205 non-null    float64
13  curb-weight            205 non-null    int64
14  engine-type            205 non-null    object
15  num-of-cylinders       205 non-null    object
16  engine-size            205 non-null    int64
17  fuel-system            205 non-null    object
18  bore                   205 non-null    object
19  stroke                 205 non-null    object
20  compression-ratio      205 non-null    float64
21  horsepower             205 non-null    object
22  peak-rpm               205 non-null    object
23  city-mpg               205 non-null    int64
24  highway-mpg            205 non-null    int64
25  price                  205 non-null    object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB
```

In [4]:

```
1 df.describe()
```

Out[4]:

	symboling	wheel- base	length	width	height	curb-weight	engine- size
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854	126.907317
std	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204	41.642693
min	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000	61.000000
25%	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000	97.000000
50%	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000	120.000000
75%	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000	141.000000
max	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000	326.000000

In [5]:

```
1 # finding the no. of rows and columns
2
3 no_of_rows = df.shape[0]
4 print("No. of rows: ", no_of_rows)
5
6 no_of_columns = df.shape[1]
7 print("No. of columns: ", no_of_columns)
```

No. of rows: 205
No. of columns: 26

In [6]:

```
1 # Checking for missing values
2
3 df.isna().sum()           # There is no missing values in this dataset
```

Out[6]:


```
symboling           0
normalized-losses   0
make                0
fuel-type           0
aspiration          0
num-of-doors        0
body-style          0
drive-wheels        0
engine-location     0
wheel-base          0
length              0
width               0
height              0
curb-weight         0
engine-type         0
num-of-cylinders    0
engine-size         0
fuel-system         0
bore                0
stroke              0
compression-ratio   0
horsepower          0
peak-rpm            0
city-mpg            0
highway-mpg         0
price               0
dtype: int64
```

In [7]:

```
1  ## Checking for outliers
2
3  plt.figure(figsize=(20,10))
4  df.boxplot()
```

Out[7]:

<AxesSubplot:>



From the boxplot we can see that there is outliers present in some of the columns

Step 4: Feature Engineering


Outliers Handling

In [8]:

```
1 sns.boxplot(x = df['wheel-base'])
```

Out[8]:

<AxesSubplot:xlabel='wheel-base'>



In [9]:

```
1 # Ouliers finding by using IQR Method
2
3 q1=df['wheel-base'].quantile(0.25)
4 q2=df['wheel-base'].quantile(0.50)
5 q3=df['wheel-base'].quantile(0.75)
6 median=df['wheel-base'].median()
7 IQR=q3-q1
8 upper_tail = q3 + 1.5 * IQR
9 lower_tail = q1 - 1.5 * IQR
10 print("Q1 :", q1)
11 print("Q2 :", q2)
12 print("Q3 :", q3)
13 print("Median :",median)
14 print("upper_tail :", upper_tail)
15 print("lower_tail :", lower_tail)
```

```
Q1 : 94.5
Q2 : 97.0
Q3 : 102.4
Median : 97.0
upper_tail : 114.25000000000001
lower_tail : 82.64999999999999
```

In [10]:

```
1 # Detecting outliers in the column
2
3 df['wheel-base'].loc[(df['wheel-base'] > upper_tail)]
```

Out[10]:

```
70    115.6
71    115.6
73    120.9
Name: wheel-base, dtype: float64
```

In [11]:

```
1 # Calculating the median of good data of Item_Visibility column
2
3 median_wheel_base = df['wheel-base'].loc[(df['wheel-base'] <= upper_tail)].median()
4 median_wheel_base
```

Out[11]:

```
96.9
```

In [12]:

```
1 # Imputing outliers by median value
2
3 df['wheel-base'].loc[(df['wheel-base'] > upper_tail)] = median_wheel_base
```

In [13]:

```
1 # Checking outliers after imputation
2
3 df[['wheel-base']].loc[(df['wheel-base'] > upper_tail)]
```

Out[13]:


wheel-base

In [14]:

```
1 sns.boxplot(x = df['width'])
```

Out[14]:

<AxesSubplot:xlabel='width'>



In [15]:

```
1 # Ouliers finding by using IQR Method
2
3 q1= df['width'].quantile(0.25)
4 q2= df['width'].quantile(0.50)
5 q3= df['width'].quantile(0.75)
6 median= df['width'].median()
7 IQR=q3-q1
8 upper_tail = q3 + 1.5 * IQR
9 lower_tail = q1 - 1.5 * IQR
10 print("Q1 :", q1)
11 print("Q2 :", q2)
12 print("Q3 :", q3)
13 print("Median :",median)
14 print("upper_tail :", upper_tail)
15 print("lower_tail :", lower_tail)
```

Q1 : 64.1

Q2 : 65.5

Q3 : 66.9

Median : 65.5

upper_tail : 71.10000000000002

lower_tail : 59.899999999999998

In [16]:

```
1 # Detecting outliers in the column
2
3 df['width'].loc[(df['width'] > upper_tail)]
```

Out[16]:

```
6      71.4
7      71.4
8      71.4
70     71.7
71     71.7
73     71.7
74     72.0
129    72.3
Name: width, dtype: float64
```

In [17]:

```
1 # Calculating the median of good data of Item_Visibility column
2
3 median_width = df['width'].loc[(df['width'] <= upper_tail)].median()
4 median_width
```

Out[17]:

65.4

In [18]:

```
1 # Imputing outliers by median value
2
3 df['width'].loc[(df['width'] > upper_tail)] = median_width
```

In [19]:

```
1 # Checking outliers after imputation
2 df[['width']].loc[(df['width'] > upper_tail)]
```

Out[19]:


width

In [20]:

```
1 sns.boxplot(x = df['engine-size'])
```

Out[20]:

<AxesSubplot:xlabel='engine-size'>



In [21]:

```
1 # Ouliers finding by using IQR Method
2
3 q1= df['engine-size'].quantile(0.25)
4 q2= df['engine-size'].quantile(0.50)
5 q3= df['engine-size'].quantile(0.75)
6 median= df['engine-size'].median()
7 IQR=q3-q1
8 upper_tail = q3 + 1.5 * IQR
9 lower_tail = q1 - 1.5 * IQR
10 print("Q1 :", q1)
11 print("Q2 :", q2)
12 print("Q3 :", q3)
13 print("Median :",median)
14 print("upper_tail :", upper_tail)
15 print("lower_tail :", lower_tail)
```

```
Q1 : 97.0
Q2 : 120.0
Q3 : 141.0
Median : 120.0
upper_tail : 207.0
lower_tail : 31.0
```

In [22]:

```
1 # Detecting outliers in the column
2
3 df['engine-size'].loc[(df['engine-size'] > upper_tail)]
```

Out[22]:

```
15    209
16    209
17    209
47    258
48    258
49    326
71    234
72    234
73    308
74    304
```

Name: engine-size, dtype: int64

In [23]:

```
1 # Imputing outliers by upper_tail value
2
3 df['engine-size'].loc[(df['engine-size'] > upper_tail)] = upper_tail
```

In [24]:

```
1 # Checking outliers after imputation
2
3 df[['engine-size']].loc[(df['engine-size'] > upper_tail)]
```

Out[24]:


engine-size

In [25]:

```
1 sns.boxplot(x = df['compression-ratio'])
```

Out[25]:

<AxesSubplot:xlabel='compression-ratio'>



In [26]:

```
1 # Ouliers finding by using IQR Method
2
3 q1= df['compression-ratio'].quantile(0.25)
4 q2= df['compression-ratio'].quantile(0.50)
5 q3= df['compression-ratio'].quantile(0.75)
6 median= df['compression-ratio'].median()
7 IQR=q3-q1
8 upper_tail = q3 + 1.5 * IQR
9 lower_tail = q1 - 1.5 * IQR
10 print("Q1 :", q1)
11 print("Q2 :", q2)
12 print("Q3 :", q3)
13 print("Median :",median)
14 print("upper_tail :", upper_tail)
15 print("lower_tail :", lower_tail)
```

Q1 : 8.6

Q2 : 9.0

Q3 : 9.4

Median : 9.0

upper_tail : 10.600000000000001

lower_tail : 7.399999999999999

In [27]:

```
1 # Detecting outliers in the column
2
3 outliers = df['compression-ratio'].loc[(df['compression-ratio'] > upper_tail) | (df[
4 outliers
```

Out[27]:

```
9      7.0
29     7.0
49    11.5
63    22.7
66    22.0
67    21.5
68    21.5
69    21.5
70    21.5
82     7.0
83     7.0
84     7.0
90    21.9
108   21.0
110   21.0
112   21.0
114   21.0
116   21.0
117     7.0
124     7.0
158   22.5
159   22.5
174   22.5
182   23.0
184   23.0
187   23.0
192   23.0
203   23.0
```

Name: compression-ratio, dtype: float64

In [28]:

```
1 # Imputing outliers by upper_tail and lower_tail values
2
3 df['compression-ratio'].loc[(df['compression-ratio'] > upper_tail)] = upper_tail
4 df['compression-ratio'].loc[(df['compression-ratio'] < lower_tail)] = lower_tail
```

In [29]:

```
1 # Checking outliers after imputation
2
3 df[['compression-ratio']].loc[(df['compression-ratio'] > upper_tail) | (df['compress:
```

Out[29]:


compression-ratio

In [30]:

```
1 sns.boxplot(x = df['city-mpg'])
```

Out[30]:

<AxesSubplot:xlabel='city-mpg'>



In [31]:

```
1 # Ouliers finding by using IQR Method
2
3 q1= df['city-mpg'].quantile(0.25)
4 q2= df['city-mpg'].quantile(0.50)
5 q3= df['city-mpg'].quantile(0.75)
6
7 IQR=q3-q1
8 upper_tail = q3 + 1.5 * IQR
9 lower_tail = q1 - 1.5 * IQR
10 print("Q1 :", q1)
11 print("Q2 :", q2)
12 print("Q3 :", q3)
13
14 print("upper_tail :", upper_tail)
15 print("lower_tail :", lower_tail)
```

```
Q1 : 19.0
Q2 : 24.0
Q3 : 30.0
upper_tail : 46.5
lower_tail : 2.5
```

In [32]:

```
1 # Detecting outliers in the column
2
3 df['city-mpg'].loc[(df['city-mpg'] > upper_tail)]
4
```

Out[32]:

```
18    47
30    49
Name: city-mpg, dtype: int64
```

In [33]:

```
1 # Imputation of outliers with upper_tail,value
2
3 df['city-mpg'].loc[(df['city-mpg'] > upper_tail)] = upper_tail
```

In [34]:

```
1 # Rechecking outliers after imputation
2
3 df[['city-mpg']].loc[(df['city-mpg'] > upper_tail)]
```

Out[34]:


city-mpg

In [35]:

```
1 sns.boxplot(x = df['highway-mpg'])
```

Out[35]:

<AxesSubplot:xlabel='highway-mpg'>



In [36]:

```
1 # Ouliers finding by using IQR Method
2
3 q1= df['highway-mpg'].quantile(0.25)
4 q2= df['highway-mpg'].quantile(0.50)
5 q3= df['highway-mpg'].quantile(0.75)
6
7 IQR=q3-q1
8 upper_tail = q3 + 1.5 * IQR
9 lower_tail = q1 - 1.5 * IQR
10 print("Q1 :", q1)
11 print("Q2 :", q2)
12 print("Q3 :", q3)
13
14 print("upper_tail :", upper_tail)
15 print("lower_tail :", lower_tail)
```

```
Q1 : 25.0
Q2 : 30.0
Q3 : 34.0
upper_tail : 47.5
lower_tail : 11.5
```

In [37]:

```
1 # Detecting outliers in the column
2
3 df['highway-mpg'].loc[(df['highway-mpg'] > upper_tail)]
```

Out[37]:

18 53

30 54

90 50

Name: highway-mpg, dtype: int64

In [38]:

```
1 # Imputation of outliers with upper_tail,value
2
3 df['highway-mpg'].loc[(df['highway-mpg'] > upper_tail)] = upper_tail
```

In [39]:

```
1 # Rechecking the outliers after imputation
2
3 df[['highway-mpg']].loc[(df['highway-mpg'] > upper_tail)]
```

Out[39]:

highway-mpg

Converting Categorical Columns into Numerical column

In [40]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   symboling              205 non-null    int64
1   normalized-losses      205 non-null    object
2   make                   205 non-null    object
3   fuel-type              205 non-null    object
4   aspiration              205 non-null    object
5   num-of-doors           205 non-null    object
6   body-style             205 non-null    object
7   drive-wheels           205 non-null    object
8   engine-location        205 non-null    object
9   wheel-base             205 non-null    float64
10  length                 205 non-null    float64
11  width                  205 non-null    float64
12  height                 205 non-null    float64
13  curb-weight            205 non-null    int64
14  engine-type            205 non-null    object
15  num-of-cylinders       205 non-null    object
16  engine-size            205 non-null    int64
17  fuel-system            205 non-null    object
18  bore                   205 non-null    object
19  stroke                 205 non-null    object
20  compression-ratio      205 non-null    float64
21  horsepower             205 non-null    object
22  peak-rpm               205 non-null    object
23  city-mpg               205 non-null    float64
24  highway-mpg            205 non-null    float64
25  price                  205 non-null    object
dtypes: float64(7), int64(3), object(16)
memory usage: 41.8+ KB
```

1. normalized-losses

In [147]:

```
1 df['normalized-losses']
```

Out[147]:

```
0      ?
1      ?
2      ?
3     164
4     164
...
200    95
201    95
202    95
203    95
204    95
```

Name: normalized-losses, Length: 205, dtype: object

In [148]:

```
1 df['normalized-losses'].value_counts()
```

Out[148]:

```
?      41
161    11
91      8
150     7
134     6
128     6
104     6
85      5
94      5
65      5
102     5
74      5
168     5
103     5
95      5
106     4
93      4
118     4
```

In [149]:

```
1 df['normalized-losses'].replace({"?": np.nan},inplace=True)
```

In [150]:

```
1 df.replace({"?": np.nan},inplace=True)
```

In [151]:

```
1 # Checking for missing value
2 df['normalized-losses'].isna().sum()
```

Out[151]:

41

In [152]:

```
1 # converting into numeric data type
2
3 df['normalized-losses'] = df['normalized-losses'].astype(float)
```

```
1 # Outliers checking
2
3 sns.boxplot(x = df['normalized-losses'])
```

In [154]:

```
1 # imputaion of missing value with median value because outliers are present in the d
2
3 df['normalized-losses'].fillna(df['normalized-losses'].median(), inplace=True)
```

In [155]:

```
1 # Rechecking for missing value
2 df['normalized-losses'].isna().sum()
```

Out[155]:

0

2. make

In [156]:

```
1 df['make'].isna().sum()
```

Out[156]:

0

In [157]:

```
1 df['make'].nunique()
```

Out[157]:

22

In [158]:

```
1 df['make'].value_counts()
```

Out[158]:

toyota	32
nissan	18
mazda	17
mitsubishi	13
honda	13
volkswagen	12
subaru	12
peugot	11
volvo	11
dodge	9
mercedes-benz	8
bmw	8
audi	7
plymouth	7
saab	6
porsche	5
isuzu	4
luxgen	3

In [159]:

```
1 # Converting categorical values into numeric value by using get_dummies
2
3 df = pd.get_dummies(df, columns = ['make'])
```

3. fuel-type

In [217]:

```
1 df['fuel-type'].value_counts()
```

Out[217]:

1	185
0	20

Name: fuel-type, dtype: int64

In [162]:

```
1 df['fuel-type'].replace({"gas":1, "diesel":0}, inplace=True)
```

In [163]:

```
1 fuel_type_dict = {"gas":1, "diesel":0}
```

In [164]:

```
1 df['fuel-type'].value_counts()
```

Out[164]:

1	185
0	20

Name: fuel-type, dtype: int64

4. aspiration

In [165]:

```
1 df['aspiration'].value_counts()
```

Out[165]:

```
std      168
turbo     37
Name: aspiration, dtype: int64
```

In [166]:

```
1 df['aspiration'].value_counts().to_dict()
```

Out[166]:

```
{'std': 168, 'turbo': 37}
```

In [167]:

```
1 df['aspiration'].replace({'std': 0, 'turbo': 1}, inplace=True)
```

In [168]:

```
1 df['aspiration'].value_counts().to_dict()
```

Out[168]:

```
{0: 168, 1: 37}
```

In [181]:

```
1 aspiration_dict = {'std': 0, 'turbo': 1}
```

5. num-of-doors

In [182]:

```
1 df['num-of-doors'].unique()
```

Out[182]:

```
array([2, 4], dtype=int64)
```

In [183]:

```
1 df['num-of-doors'].isna().sum()
```

Out[183]:

```
0
```

In [184]:

```
1 df['num-of-doors'].value_counts()
```

Out[184]:

```
4    116
2     89
Name: num-of-doors, dtype: int64
```

In [185]:

```
1 # Imputation of missing value with mode
2
3 df['num-of-doors'].fillna(df['num-of-doors'].mode()[0], inplace=True)
```

In [186]:

```
1 df['num-of-doors'].value_counts().to_dict()
```

Out[186]:

```
{4: 116, 2: 89}
```

In [187]:

```
1 df['num-of-doors'].replace({'four': 4, 'two': 2}, inplace=True)
```

In [188]:

```
1 df['num-of-doors'].value_counts().to_dict()
```

Out[188]:

```
{4: 116, 2: 89}
```

In [189]:

```
1 num_of_doors_dict = {'four': 4, 'two': 2}
```

6. body-style

In [178]:

```
1 df['body-style'].unique()
```

Out[178]:

```
array(['convertible', 'hatchback', 'sedan', 'wagon', 'hardtop'],
      dtype=object)
```


In [179]:

```
1 df['body-style'].value_counts()
```

Out[179]:

```
sedan      96
hatchback  70
wagon      25
hardtop     8
convertible 6
Name: body-style, dtype: int64
```

In [180]:

```
1 df = pd.get_dummies(df, columns=['body-style'])
```

7. drive-wheels

In [212]:

```
1 df['drive-wheels'].unique()
```

Out[212]:

```
array(['rwd', 'fwd', '4wd'], dtype=object)
```

In [213]:

```
1 df['drive-wheels'].value_counts()
```

Out[213]:

```
fwd    120
rwd     76
4wd      9
Name: drive-wheels, dtype: int64
```

In [214]:

```
1 df['drive-wheels'].replace({'fwd': 0, 'rwd': 1, '4wd': 2}, inplace = True)
```

In [215]:

```
1 drive_wheels_dict = {'fwd': 0, 'rwd': 1, '4wd': 2}
```

In [216]:

```
1 df['drive-wheels'].value_counts()
```

Out[216]:

```
0    120
1     76
2      9
Name: drive-wheels, dtype: int64
```

8. engine-location

In [208]:

```
1 df['engine-location'].unique()
```

Out[208]:

```
array(['front', 'rear'], dtype=object)
```

In [209]:

```
1 df['engine-location'].replace({'front': 0, 'rear': 1}, inplace=True)
```

In [301]:

```
1 engine_location_dict = {'front': 0, 'rear': 1}
```

In [211]:

```
1 df['engine-location'].value_counts()
```

Out[211]:

```
0    202
1      3
Name: engine-location, dtype: int64
```

9. engine-type

In [227]:

```
1 df['engine-type'].unique()
```

Out[227]:

```
array(['dohc', 'ohcv', 'ohc', 'l', 'rotor', 'ohcf', 'dohcv'], dtype=object)
```

In [228]:

```
1 df['engine-type'].value_counts()
```

Out[228]:

```
ohc      148
ohcf     15
ohcv     13
dohc     12
l         12
rotor      4
dohcv      1
Name: engine-type, dtype: int64
```

In [229]:

```
1 df = pd.get_dummies(df, columns=['engine-type'])
2 df
```

Out[229]:

	symboling	normalized- losses	fuel- type	aspiration	num- of- doors	drive- wheels	engine- location	wheel- base	length	width
0	3	115.0	1	0	2	1	0	88.6	168.8	64.1
1	3	115.0	1	0	2	1	0	88.6	168.8	64.1
2	1	115.0	1	0	2	1	0	94.5	171.2	65.5
3	2	164.0	1	0	4	0	0	99.8	176.6	66.2
4	2	164.0	1	0	4	2	0	99.4	176.6	66.4
...
200	-1	95.0	1	0	4	1	0	109.1	188.8	68.9
201	-1	95.0	1	1	4	1	0	109.1	188.8	68.8
202	-1	95.0	1	0	4	1	0	109.1	188.8	68.9
203	-1	95.0	0	1	4	1	0	109.1	188.8	68.9
204	-1	95.0	1	1	4	1	0	109.1	188.8	68.9

205 rows × 64 columns

10. num-of-cylinders

In [230]:

```
1 df['num-of-cylinders'].unique()
```

Out[230]:

```
array(['four', 'six', 'five', 'three', 'twelve', 'two', 'eight'],
      dtype=object)
```

In [231]:

```
1 df['num-of-cylinders'].value_counts().to_dict()
```

Out[231]:

```
{'four': 159,
 'six': 24,
 'five': 11,
 'eight': 5,
 'two': 4,
 'three': 1,
 'twelve': 1}
```

In [232]:

```
1 df['num-of-cylinders'].replace({'four': 4,  
2 'six': 6,  
3 'five': 5,  
4 'eight': 8,  
5 'two': 2,  
6 'three': 3,  
7 'twelve': 12},inplace=True)
```

In [302]:

```
1 num_of_cylinders_dict = {'four': 4,  
2 'six': 6,  
3 'five': 5,  
4 'eight': 8,  
5 'two': 2,  
6 'three': 3,  
7 'twelve': 12}
```

In [234]:

```
1 df['num-of-cylinders'].value_counts()
```

Out[234]:

```
4      159  
6       24  
5       11  
8        5  
2        4  
3        1  
12       1  
Name: num-of-cylinders, dtype: int64
```

11. fuel-system

In [223]:

```
1 df['fuel-system'].unique()
```

Out[223]:

```
array(['mpfi', '2bbl', 'mfi', '1bbl', 'spfi', '4bbl', 'idi', 'spdi'],  
      dtype=object)
```

In [224]:

```
1 df['fuel-system'].value_counts()
```

Out[224]:

```
mpfi      94
2bbl      66
idi       20
1bbl      11
spdi       9
4bbl       3
mfi        1
spfi        1
Name: fuel-system, dtype: int64
```


In [225]:

```
1 df = pd.get_dummies(df, columns=['fuel-system'])
2 df
```

Out[225]:

	symboling	normalized-losses	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	wheel-base	length	width
0	3	115.0	1	0	2	1	0	88.6	168.8	64.1
1	3	115.0	1	0	2	1	0	88.6	168.8	64.1
2	1	115.0	1	0	2	1	0	94.5	171.2	65.5
3	2	164.0	1	0	4	0	0	99.8	176.6	66.2
4	2	164.0	1	0	4	2	0	99.4	176.6	66.4
...
200	-1	95.0	1	0	4	1	0	109.1	188.8	68.9
201	-1	95.0	1	1	4	1	0	109.1	188.8	68.8
202	-1	95.0	1	0	4	1	0	109.1	188.8	68.9
203	-1	95.0	0	1	4	1	0	109.1	188.8	68.9
204	-1	95.0	1	1	4	1	0	109.1	188.8	68.9

205 rows × 58 columns



12. bore

In [219]:

```
1 df['bore'].unique()
```

Out[219]:

```
array(['3.47', '2.68', '3.19', '3.13', '3.5', '3.31', '3.62', '2.91',  
      '3.03', '2.97', '3.34', '3.6', '2.92', '3.15', '3.43', '3.63',  
      '3.54', '3.08', nan, '3.39', '3.76', '3.58', '3.46', '3.8', '3.78',  
      '3.17', '3.35', '3.59', '2.99', '3.33', '3.7', '3.61', '3.94',  
      '3.74', '2.54', '3.05', '3.27', '3.24', '3.01'], dtype=object)
```

In [220]:

```
1 df['bore'].isna().sum()
```

Out[220]:

4

In [221]:

```
1 # Imputation of missing value with median  
2  
3 df['bore'] = df['bore'].fillna(df['bore'].median()).astype(float)
```

In [222]:

```
1 df['bore'].isna().sum()
```

Out[222]:

0

13. stroke

In [190]:

```
1 df['stroke'].unique()
```

Out[190]:

```
array(['2.68', '3.47', '3.4', '2.8', '3.19', '3.39', '3.03', '3.11',  
      '3.23', '3.46', '3.9', '3.41', '3.07', '3.58', '4.17', '2.76',  
      '3.15', nan, '3.16', '3.64', '3.1', '3.35', '3.12', '3.86', '3.29',  
      '3.27', '3.52', '2.19', '3.21', '2.9', '2.07', '2.36', '2.64',  
      '3.08', '3.5', '3.54', '2.87'], dtype=object)
```

In [191]:

```
1 df['stroke'].isna().sum()
```

Out[191]:

4

In [192]:

```
1 # Imputation of missing value
2
3 df['stroke'] = df['stroke'].fillna(df['stroke'].median()).astype(float)
```

In [193]:

```
1 df['stroke'].isna().sum()
```

Out[193]:

0

14. horsepower

In [194]:

```
1 df['horsepower'].unique()
```

Out[194]:

```
array(['111', '154', '102', '115', '110', '140', '160', '101', '121',
      '182', '48', '70', '68', '88', '145', '58', '76', '60', '86',
      '100', '78', '90', '176', '262', '135', '84', '64', '120', '72',
      '123', '155', '184', '175', '116', '69', '55', '97', '152', '200',
      '95', '142', '143', '207', '288', nan, '73', '82', '94', '62',
      '56', '112', '92', '161', '156', '52', '85', '114', '162', '134',
      '106'], dtype=object)
```

In [195]:

```
1 df['horsepower'].isna().sum()
```

Out[195]:

2

In [196]:

```
1 # Imputation of missing value
2
3 df['horsepower'] = df['horsepower'].fillna(df['horsepower'].median()).astype(float)
```

In [197]:

```
1 df['horsepower'].isna().sum()
```

Out[197]:

0

15. peak-rpm

In [198]:

```
1 df['peak-rpm'].unique()
```

Out[198]:

```
array(['5000', '5500', '5800', '4250', '5400', '5100', '4800', '6000',  
      '4750', '4650', '4200', '4350', '4500', '5200', '4150', '5600',  
      '5900', '5750', nan, '5250', '4900', '4400', '6600', '5300'],  
      dtype=object)
```

In [199]:

```
1 df['peak-rpm'].isna().sum()
```

Out[199]:

```
2
```

In [200]:

```
1 df['peak-rpm'] = df['peak-rpm'].fillna(df['peak-rpm'].median()).astype(float)
```

In [201]:

```
1 df['peak-rpm'].isna().sum()
```

Out[201]:

```
0
```

16. price

In [202]:

```
1 df['price'].isna().sum()
```

Out[202]:

```
4
```

In [203]:

```
1 df['price'] = df['price'].fillna(df['price'].median()).astype(float)
```

In [204]:

```
1 df['price'].isna().sum()
```

Out[204]:

```
0
```


In [235]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 205 entries, 0 to 204
```

```
Data columns (total 64 columns):
```

#	Column	Non-Null Count	Dtype
0	symboling	205 non-null	int64
1	normalized-losses	205 non-null	float64
2	fuel-type	205 non-null	int64
3	aspiration	205 non-null	int64
4	num-of-doors	205 non-null	int64
5	drive-wheels	205 non-null	int64
6	engine-location	205 non-null	int64
7	wheel-base	205 non-null	float64
8	length	205 non-null	float64
9	width	205 non-null	float64
10	height	205 non-null	float64
11	curb-weight	205 non-null	int64
12	num-of-cylinders	205 non-null	int64
13	engine-size	205 non-null	int64
14	bore	205 non-null	float64
15	stroke	205 non-null	float64
16	compression-ratio	205 non-null	float64
17	horsepower	205 non-null	float64
18	peak-rpm	205 non-null	float64
19	city-mpg	205 non-null	int64
20	highway-mpg	205 non-null	int64
21	price	205 non-null	float64
22	make_alfa-romero	205 non-null	uint8
23	make_audi	205 non-null	uint8
24	make_bmw	205 non-null	uint8
25	make_chevrolet	205 non-null	uint8
26	make_dodge	205 non-null	uint8
27	make_honda	205 non-null	uint8
28	make_isuzu	205 non-null	uint8
29	make_jaguar	205 non-null	uint8
30	make_mazda	205 non-null	uint8
31	make_mercedes-benz	205 non-null	uint8
32	make_mercury	205 non-null	uint8
33	make_mitsubishi	205 non-null	uint8
34	make_nissan	205 non-null	uint8
35	make_peugot	205 non-null	uint8
36	make_plymouth	205 non-null	uint8
37	make_porsche	205 non-null	uint8
38	make_renault	205 non-null	uint8
39	make_saab	205 non-null	uint8
40	make_subaru	205 non-null	uint8
41	make_toyota	205 non-null	uint8
42	make_volkswagen	205 non-null	uint8
43	make_volvo	205 non-null	uint8
44	body-style_convertible	205 non-null	uint8
45	body-style_hardtop	205 non-null	uint8
46	body-style_hatchback	205 non-null	uint8
47	body-style_sedan	205 non-null	uint8
48	body-style_wagon	205 non-null	uint8
49	fuel-system_1bbl	205 non-null	uint8
50	fuel-system_2bbl	205 non-null	uint8
51	fuel-system_4bbl	205 non-null	uint8
52	fuel-system_idi	205 non-null	uint8
53	fuel-system_mfi	205 non-null	uint8
54	fuel-system_mphi	205 non-null	uint8
55	fuel-system_spdi	205 non-null	uint8

```
56 fuel-system_spfi      205 non-null    uint8
57 engine-type_dohc      205 non-null    uint8
58 engine-type_dohcv     205 non-null    uint8
59 engine-type_l         205 non-null    uint8
60 engine-type_ohc       205 non-null    uint8
61 engine-type_ohcf      205 non-null    uint8
62 engine-type_ohcv      205 non-null    uint8
63 engine-type_rotor     205 non-null    uint8
dtypes: float64(11), int64(11), uint8(42)
memory usage: 43.8 KB
```

Step 5: Feature Selection

In [236]:

```
1 df.corr()
```

Out[236]:

	symboling	normalized-losses	fuel-type	aspiration	num-of-doors	drive-wheels	engine-location	
symboling	1.000000	0.457484	0.194311	-0.059866	-0.663595	-0.111150	0.212471	-0
normalized-losses	0.457484	1.000000	0.104668	-0.011273	-0.348850	0.133824	-0.021510	-0
fuel-type	0.194311	0.104668	1.000000	-0.401397	-0.188496	-0.051874	0.040070	-0
aspiration	-0.059866	-0.011273	-0.401397	1.000000	0.052803	0.153897	-0.057191	0
num-of-doors	-0.663595	-0.348850	-0.188496	0.052803	1.000000	-0.003230	-0.139129	0
...
engine-type_l	-0.133979	0.170806	-0.268163	0.207156	0.176489	0.197050	-0.030388	0
engine-type_ohc	-0.082855	-0.156069	-0.020584	-0.020162	0.027539	-0.429386	-0.196371	-0
engine-type_ohcf	0.037513	-0.210771	0.092384	-0.034450	0.019357	0.197807	0.433727	-0
engine-type_ohcv	-0.013597	0.130717	0.085556	-0.070070	-0.054764	0.139453	-0.031711	0
engine-type_rotor	0.245950	0.130721	0.046383	-0.066203	-0.161052	0.131758	-0.017192	-0

64 rows × 64 columns



In [237]:

```

1  # Checking for Multicollinearity
2
3  df1 = df.drop('price', axis=1)    # Dropping dependant variable from dataset
4
5  vif_list = []
6
7  for i in range(df1.shape[1]):
8      vif = variance_inflation_factor(df1,i)
9      vif_list.append(vif)
10
11 s1 = pd.Series(vif_list, index=df1.columns)
12 s1
13
14 # s1.sort_values().plot(kind = 'barh')

```

Out[237]:

```

symboling          6.062736
normalized-losses  3.112823
fuel-type          inf
aspiration         6.187690
num-of-doors       3.817695
...
engine-type_l      inf
engine-type_ohc    inf
engine-type_ohcf   inf
engine-type_ohcv   inf
engine-type_rotor  inf
Length: 63, dtype: float64

```

Train-Test Split

In [238]:

```

1  x = df.drop('price', axis= 1) # independent variables
2  y = df['price'] # dependent variables

```

In [251]:

```

1  x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=0.2, random_state=42)
2
3  print("Training data counts", x_train.shape, y_train.shape)
4  print("Testing data counts", x_test.shape, y_test.shape)

```

Training data counts (164, 63) (164,)

Testing data counts (41, 63) (41,)

Model Training

In [258]:

```
1 model_linear = LinearRegression()  
2  
3 model_linear.fit(x_train, y_train)
```

Out[258]:

LinearRegression()

In [260]:

```
1 # Model Evaluation for Training Data  
2  
3 y_pred_train = model_linear.predict(x_train)  
4  
5 mse = mean_squared_error(y_train, y_pred_train)  
6 print("Mean Sqaured Error :",mse)  
7  
8 rmse = np.sqrt(mse)  
9 print("Root Mean Sqaured Error :",rmse)  
10  
11 mae = mean_absolute_error(y_train, y_pred_train)  
12 print("Mean Absolute Error :",mae)  
13  
14 r_squared_value = r2_score(y_train, y_pred_train)  
15 print("R Squared Value :",r_squared_value)  
16  
17 adj_r2 = 1 - (((1 - r_squared_value) * (x_train.shape[0] - 1)) / (x_train.shape[0]  
18 print("Adjusted R Squared Value :",adj_r2)
```

Mean Sqaured Error : 2391392.9332026816

Root Mean Sqaured Error : 1546.4129245459253

Mean Absolute Error : 1129.2985382288748

R Squared Value : 0.962437619451845

Adjusted R Squared Value : 0.9387733197065073

In [261]:

```

1 # Model Evaluation for Testing Data
2 y_pred = model_linear.predict(x_test)
3
4 mse = mean_squared_error(y_test, y_pred)
5 print("Mean Sqaured Error :",mse)
6
7 rmse = np.sqrt(mse)
8 print("Root Mean Sqaured Error :",rmse)
9
10 mae = mean_absolute_error(y_test, y_pred)
11 print("Mean Absolute Error :",mae)
12
13 r_squared_value = r2_score(y_test, y_pred)
14 print("R Squared Value :",r_squared_value)
15
16 adj_r2 = 1 - (((1 - r_squared_value) * (x_test.shape[0] - 1)) / (x_test.shape[0] -
17 print("Adjusted R Squared Value :",adj_r2)

```

Mean Sqaured Error : 6792666.853525117

Root Mean Sqaured Error : 2606.2745161485036

Mean Absolute Error : 1592.968674573721

R Squared Value : 0.8701141652233484

Adjusted R Squared Value : 1.2258884083072201

In [262]:


```

1 s2 = pd.Series(model_linear.coef_, index=x.columns)
2 s2.sort_values().plot(kind = "barh")

```

Out[262]:

<AxesSubplot:>



```
1 By comparing the accuracy of training and testing data. It is observed that model  
is overfitted therefore we are checking the assumptions of linear regression.
```

In [263]:

```
1 Residual = y_train - y_pred_train  
2 Residual
```

Out[263]:

```
111    1805.501237  
93     636.621267  
148   -1326.147280  
21    -818.147122  
28     771.501224  
  
...  
56    -888.691189  
182    -22.879388  
204    1489.372537  
92     431.285687  
126   -1131.585946  
Name: price, Length: 164, dtype: float64
```


1. Assumption of Normality of Residual

In [264]:

```
1 # KDE plot  
2  
3 sns.kdeplot(Residual, fill= True )
```

Out[264]:

<AxesSubplot:xlabel='price', ylabel='Density'>



In [265]:

```
1 # Hypothesis Testing
2
3 _ , p_val = shapiro(Residual)
4
5 print("P_Value:",p_val)
6
7 if p_val >= 0.05:
8     print("Null Hypothesis is Accepted")
9     print("--> Data is Normally Distributed")
10
11 else:
12     print("Null Hypothesis is Rejected >> Alternate Hypothesis is Accepted")
13     print("Data is NOT Normally Distributed")
```

P_Value: 7.378377631539479e-06


Null Hypothesis is Rejected >> Alternate Hypothesis is Accepted

Data is NOT Normally Distributed

In [266]:

```
1 # By Using QQ plot (Visualization Method)
2
3 sm.qqplot(Residual, line = '45', fit=True)
4
5 # If 90% of datapoints are on the line --> data is Normally distributed
```


Out[266]:



```
1 | sns.scatterplot(x = y_train, y = Residual)
```

Out[267]:

<AxesSubplot:xlabel='price', ylabel='price'>



Model Building By using Regularization Techniques

1. Ridge Regression

In [268]:

```
1 | ridge_reg_model = Ridge()           # alpha= 1.0 (bydefault)
2 |
3 | ridge_reg_model.fit(x_train, y_train)
```

Out[268]:

Ridge()

In [269]:

```
1 # Model Evaluation on Training Dataset
2
3 y_pred_train = ridge_reg_model.predict(x_train)
4
5 mse = mean_squared_error(y_train, y_pred_train)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_train, y_pred_train)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_train, y_pred_train)
15 print("R2 Scored :", r2)
```

MSE : 3382887.8923862255
RMSE : 1839.262866581671
MAE : 1294.086249500661
R2 Scored : 0.9468638881543489

In [270]:

```
1 # Model Evaluation on Testing Dataset
2
3 y_pred = ridge_reg_model.predict(x_test)
4
5 mse = mean_squared_error(y_test, y_pred)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_test, y_pred)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_test, y_pred)
15 print("R2 Scored :", r2)
```

MSE : 6025948.636963595
RMSE : 2454.7807716705775
MAE : 1678.821286187017
R2 Scored : 0.8847749512951515

In [271]:

```
1 s3 = pd.Series(ridge_reg_model.coef_, index=x.columns)
2 s3.sort_values()
```

Out[271]:

```
engine-type_dohcv    -6004.355103
bore                 -4197.134277
make_subaru          -2842.436517
make_mitsubishi      -2661.321138
make_plymouth        -2289.347167
...
aspiration           3142.450446
engine-type_rotor    3330.204138
make_mercedes-benz   4130.231640
engine-location      5637.771822
make_bmw             6226.981653
Length: 63, dtype: float64
```

2. Lasso Regression

In [272]:

```
1 lasso_reg_model = Lasso()                # alpha= 1.0 (bydefault)
2
3 lasso_reg_model.fit(x_train, y_train)
```

Out[272]:

Lasso()

In [273]:

```
1 # Model Evaluation on Training Dataset
2
3 y_pred_train = lasso_reg_model.predict(x_train)
4
5 mse = mean_squared_error(y_train, y_pred_train)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_train, y_pred_train)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_train, y_pred_train)
15 print("R2 Scored :", r2)
```

```
MSE : 2404575.9929006225
RMSE : 1550.6695305256444
MAE : 1134.8654267419054
R2 Scored : 0.9622305488787544
```

In [274]:

```
1 # Model Evaluation on Testing Dataset
2
3 y_pred = lasso_reg_model.predict(x_test)
4
5 mse = mean_squared_error(y_test, y_pred)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_test, y_pred)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_test, y_pred)
15 print("R2 Scored :", r2)
```

MSE : 6455877.75814986
RMSE : 2540.841938836389
MAE : 1550.8424925158415
R2 Scored : 0.876554070747905

In [275]:

```
1 s4 = pd.Series(lasso_reg_model.coef_, index=x.columns)
2 s4.sort_values()
```

Out[275]:

engine-type_dohcv	-20959.411383
make_peugot	-6490.384046
bore	-3728.854131
make_mitsubishi	-3265.543817
make_plymouth	-2947.792748
...	
engine-location	7716.240019
engine-type_rotor	7920.855914
make_porsche	7934.902085
make_bmw	8380.363052
fuel-system_idi	9545.428879

Length: 63, dtype: float64

3. Hyperparameter Tuning

A. Ridge Regression

BY Using GridsearchCV

In [276]:

```
1 # Model instance
2 ridge_reg_model = Ridge()
3
4 # Defined param_grid
5 param_grid = {"alpha": np.arange(0.01,3,0.01)}
6
7 gscv_ridge_model = GridSearchCV(ridge_reg_model, param_grid, n_jobs=-1)
8
9 gscv_ridge_model.fit(x_train, y_train)
10
11 gscv_ridge_model.best_estimator_
```

Out[276]:

Ridge(alpha=0.01)

In [277]:

```
1 # Rebuild the model by using new alpha value
2
3 ridge_reg_model = Ridge(alpha = 0.01)
4
5 ridge_reg_model.fit(x_train, y_train)
```

Out[277]:

Ridge(alpha=0.01)

In [278]:

```
1 # Evaluation Matrix on Training Dataset
2
3 y_pred_train = ridge_reg_model.predict(x_train)
4
5 mse = mean_squared_error(y_train, y_pred_train)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_train, y_pred_train)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_train, y_pred_train)
15 print("R2 Scored :", r2)
```

MSE : 2393827.1930172667

RMSE : 1547.199790918182

MAE : 1132.2192184501341

R2 Scored : 0.962399383747357

In [279]:

```
1 # Evaluation Matrix on Testing Dataset
2
3 y_pred = ridge_reg_model.predict(x_test)
4
5 mse = mean_squared_error(y_test, y_pred)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_test, y_pred)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_test, y_pred)
15 print("R2 Scored :", r2)
```

MSE : 6677560.767209952
RMSE : 2584.097669827894
MAE : 1581.9639298494208
R2 Scored : 0.8723151638048054

By using Randomized Search CV

In [280]:

```
1 # Model instance
2 ridge_reg_model = Ridge()
3
4 # Defined param_grid
5 param_grid = {"alpha": np.arange(0.01,3,0.01)}
6
7 rscv_ridge_model = RandomizedSearchCV(ridge_reg_model, param_grid, n_jobs=-1)
8
9 rscv_ridge_model.fit(x_train, y_train)
10
11 rscv_ridge_model.best_estimator_
```

Out[280]:

Ridge(alpha=0.09999999999999999)

In [281]:

```
1 # Rebuild the model by using new alpha value
2
3 ridge_reg_model = Ridge(alpha = 0.09999999999999999)
4
5 ridge_reg_model.fit(x_train, y_train)
```

Out[281]:

Ridge(alpha=0.09999999999999999)

In [282]:

```
1 # Model Evaluation on Training Dataset
2
3 y_pred_train = ridge_reg_model.predict(x_train)
4
5 mse = mean_squared_error(y_train, y_pred_train)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_train, y_pred_train)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_train, y_pred_train)
15 print("R2 Scored :", r2)
```

MSE : 2488819.012457033
RMSE : 1577.5991292014055
MAE : 1151.6574760664
R2 Scored : 0.960907316583815

In [283]:

```
1 # Model Evaluation on Testing Dataset
2
3 y_pred = ridge_reg_model.predict(x_test)
4
5 mse = mean_squared_error(y_test, y_pred)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_test, y_pred)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_test, y_pred)
15 print("R2 Scored :", r2)
```

MSE : 6262444.761345223
RMSE : 2502.4877145243336
MAE : 1579.765486212505
R2 Scored : 0.880252795682469

B. Lasso Regression

By using GridSearchCV

In [285]:

```
1 # Model instance
2 lasso_reg_model = Lasso()
3
4 # Defined param_grid
5 param_grid = {"alpha": np.arange(0.01,3,0.01)}
6
7
8 gscv_lasso_model = GridSearchCV(lasso_reg_model, param_grid, n_jobs=-1, cv= 5)
9
10 gscv_lasso_model.fit(x_train, y_train)
11 gscv_lasso_model.best_estimator_
```

Out[285]:

Lasso(alpha=0.13)

In [286]:

```
1 # Rebuild model by using new value of alpha
2
3 lasso_reg_model = Lasso(alpha = 0.13)
4 lasso_reg_model.fit(x_train, y_train)
```

Out[286]:

Lasso(alpha=0.13)

In [287]:

```
1 # Model Evaluation on Training Dataset
2
3 y_pred_train = lasso_reg_model.predict(x_train)
4
5 mse = mean_squared_error(y_train, y_pred_train)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_train, y_pred_train)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_train, y_pred_train)
15 print("R2 Scored :", r2)
```

MSE : 2391726.5346919987
RMSE : 1546.5207837892121
MAE : 1129.8624048067356
R2 Scored : 0.962432379465593

In [288]:

```
1 # Model Evaluation on Testing Dataset
2
3 y_pred = lasso_reg_model.predict(x_test)
4
5 mse = mean_squared_error(y_test, y_pred)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_test, y_pred)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_test, y_pred)
15 print("R2 Scored :", r2)
```

MSE : 6733725.918703463
RMSE : 2594.942372906085
MAE : 1585.3126714642851
R2 Scored : 0.8712412030550145

By using RandomizesSearchCV

In [289]:

```
1 # Model instance
2 lasso_model = Lasso()
3
4 # Defined param_grid
5 param_grid = {"alpha": np.arange(0.01,3,0.01)}
6
7
8 rscv_lasso_model = RandomizedSearchCV(lasso_model, param_grid, n_jobs=-1)
9
10 rscv_lasso_model.fit(x_train, y_train)
11
12 rscv_lasso_model.best_estimator_
```

Out[289]:

Lasso(alpha=1.47)

In [290]:

```
1 # Rebuild the model by using new value of alpha
2
3 lasso_reg_model = Lasso(alpha=1.47)
4 lasso_reg_model.fit(x_train, y_train)
```

Out[290]:

Lasso(alpha=1.47)

In [291]:

```
1 # Model Evaluation on Training Dataset
2
3 y_pred_train = lasso_reg_model.predict(x_train)
4
5 mse = mean_squared_error(y_train, y_pred_train)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_train, y_pred_train)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_train, y_pred_train)
15 print("R2 Scored :", r2)
```

MSE : 2418392.208075672
RMSE : 1555.118068853832
MAE : 1138.2441936562661
R2 Scored : 0.962013533128254

In [292]:

```
1 # Model Evalution on Testing Dataset
2
3 y_pred = lasso_reg_model.predict(x_test)
4
5 mse = mean_squared_error(y_test, y_pred)
6 print("MSE :",mse)
7
8 rmse = np.sqrt(mse)
9 print("RMSE :",rmse)
10
11 mae = mean_absolute_error(y_test, y_pred)
12 print("MAE :",mae)
13
14 r2 = r2_score(y_test, y_pred)
15 print("R2 Scored :", r2)
```

MSE : 6329397.040614621
RMSE : 2515.8292948081
MAE : 1534.5770906645973
R2 Scored : 0.8789725690983267

- 1 By comparing all model accuracy it is observed that we get good accuracy on ridge regression model. Therefore we are creating pickle file of ridge regression model.

Creating Pickle File

In [293]:

```
1 with open("Ridge Model.pkl", "wb") as f:
2     pickle.dump(ridge_reg_model, f)
```

Creating JSON File

In [295]:

```
1 column_names = x.columns
```

In [303]:

```
1 json_data = {"fuel_type":fuel_type_dict, "aspiration": aspiration_dict, "num_of_doors": num_of_doors_dict,
2             "drive_wheels": drive_wheels_dict, "engine_location" : engine_location_dict,
3             "num_of_cylinders":num_of_cylinders_dict, "columns": list(column_names)}
4 json_data
```

Out[303]:

```
In [305]:
```

Testing for user input value

```
'compression-ratio',
```

```
24 # one-hot encoded columns
```

```

'fuel-system_2bbl',
In [307]: fuel_system_4bbl',
1 fuel-system_idi',
array = np.zeros(len(x.columns), dtype = int)
2 fuel-system_mfi',
array
3 fuel-system_mphi',
Out[307]: system_spdi',
'fuel-system_spfi',
array([0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
'engine-type_dohc', 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
'engine-type_dohc', 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
'engine-type_10', 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
'engine-type_ohc',
In [309]: 'engine-type_ohcf',
'engine-type_ohcv',
1 engine_type_col = "body-style_" + body_style
2 engine_type_col = "engine-type_" + engine_type
3 fuel_system_col = "fuel-system_" + fuel_system
4 make_col = "make_" + make

```

In [313]:

```

1 # Find the index of this column -->
2
3 body_style_index = json_data['columns'].index(body_style_col)
4 engine_type_index = json_data['columns'].index(engine_type_col)
5 fuel_system_index = json_data['columns'].index(fuel_system_col)
6 make_index = json_data['columns'].index(make_col)

```


In [316]:

```

1 array[0] = symboling
2 array[1] = normalized_losses
3 array[2] = fuel_type_dict[fuel_type]
4 array[3] = aspiration_dict[aspiration]
5 array[4] = num_of_doors_dict[num_of_doors]
6 array[5] = drive_wheels_dict[drive_wheels]
7 array[6] = engine_location_dict[engine_location]
8 array[7] = wheel_base
9 array[8] = length
10 array[9] = width
11 array[10] = height
12 array[11] = curb_weight
13 array[12] = num_of_cylinders_dict[num_of_cylinders]
14 array[13] = engine_size
15 array[14] = bore
16 array[15] = stroke
17 array[16] = compression_ratio
18 array[17] = horsepower
19 array[18] = peak_rpm
20 array[19] = city_mpg
21 array[20] = highway_mpg
22
23 array[body_style_index] = 1
24 array[engine_type_index] = 1
25 array[fuel_system_index] = 1
26 array[make_index] = 1
27 array

```

Out[316]:

```

array([[ 3, 118, 1, 1, 4, 2, 1, 88, 170, 64, 50,
        2600, 5, 130, 3, 2, 9, 111, 6000, 21, 27, 0,
         1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
         0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0,
         0, 0, 0, 0, 1, 0, 0, 0]])

```

In [317]:

```

1 prediction = ridge_reg_model.predict([array])[0] # 2D array
2 print("Prediction of your car Price is : $", round(prediction, 2))

```

Prediction of your car Price is : \$ 28388.7